

customer-churn-IBM-dataset(Predictive Model)

April 5, 2025

1 Introduction

2 Telco Customer Churn Analysis: IBM Dataset

This project analyzes customer churn data from a fictional telecommunications company that provided home phone and Internet services to 7,043 customers in California during Q3. Also this project creates a predictive model which predicts chances of customer churn. Conducted in a Jupyter Notebook, the project involves scoping, cleaning, analyzing, and visualizing the data to uncover underlying trends and then using it for Predictive ML model.

Through exploratory data analysis—using correlation matrices, chi-square tests, and other statistical methods—we seek to answer key questions, including:

- What is the most popular reason for customers canceling their subscription?
- When do cancellations most commonly occur?
- Are the correlations between these factors statistically significant?
- How do monthly charges affect the overall duration of subscriptions?
- How does the type of subscription contract influence churn? - Whats is the most common month to cancel the subscription?

Additional questions and insights will also be addressed as part of this comprehensive analysis.

Data sources:

Telco_customer_churn.xlsx was provided by **IBM**. This dataset is detailed in: <https://community.ibm.com/community/user/businessanalytics/blogs/steven-macko/2019/07/11/telco-customer-churn-1113>

Downloaded from: <https://community.ibm.com/accelerators/?context=analytics&query=telco%20churn&type=D>

The data for this project is *inspired* by real data.

2.1 Import Python Modules

Here are the primary modules that will be used in this project:

```
[581]: import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import seaborn as sns
import plotly.express as px
import scipy.stats as stats
```

```
import statsmodels.api as sm
%matplotlib inline
```

2.2 Getting data to know

Data Description 7043 observations with 33 variables

CustomerID: A unique ID that identifies each customer.

Count: A value used in reporting/dashboarding to sum up the number of customers in a filtered set.

Country: The country of the customer's primary residence.

State: The state of the customer's primary residence.

City: The city of the customer's primary residence.

Zip Code: The zip code of the customer's primary residence.

Lat Long: The combined latitude and longitude of the customer's primary residence.

Latitude: The latitude of the customer's primary residence.

Longitude: The longitude of the customer's primary residence.

Gender: The customer's gender: Male, Female

Senior Citizen: Indicates if the customer is 65 or older: Yes, No

Partner: Indicate if the customer has a partner: Yes, No

Dependents: Indicates if the customer lives with any dependents: Yes, No. Dependents could be children, parents, grandparents, etc.

Tenure Months: Indicates the total amount of months that the customer has been with the company by the end of the quarter specified above.

Phone Service: Indicates if the customer subscribes to home phone service with the company: Yes, No

Multiple Lines: Indicates if the customer subscribes to multiple telephone lines with the company: Yes, No

Internet Service: Indicates if the customer subscribes to Internet service with the company: No, DSL, Fiber Optic, Cable.

Online Security: Indicates if the customer subscribes to an additional online security service provided by the company: Yes, No

Online Backup: Indicates if the customer subscribes to an additional online backup service provided by the company: Yes, No

Device Protection: Indicates if the customer subscribes to an additional device protection plan for their Internet equipment provided by the company: Yes, No

Tech Support: Indicates if the customer subscribes to an additional technical support plan from the company with reduced wait times: Yes, No

Streaming TV: Indicates if the customer uses their Internet service to stream television programming from a third party provider: Yes, No. The company does not charge an additional fee for this service.

Streaming Movies: Indicates if the customer uses their Internet service to stream movies from a third party provider: Yes, No. The company does not charge an additional fee for this service.

Contract: Indicates the customer's current contract type: Month-to-Month, One Year, Two Year.

Paperless Billing: Indicates if the customer has chosen paperless billing: Yes, No

Payment Method: Indicates how the customer pays their bill: Bank Withdrawal, Credit Card, Mailed Check

Monthly Charge: Indicates the customer's current total monthly charge for all their services from the company.

Total Charges: Indicates the customer's total charges, calculated to the end of the quarter specified above.

Churn Label: Yes = the customer left the company this quarter. No = the customer remained with the company. Directly related to Churn Value.

Churn Value: 1 = the customer left the company this quarter. 0 = the customer remained with the company. Directly related to Churn Label.

Churn Score: A value from 0-100 that is calculated using the predictive tool IBM SPSS Modeler. The model incorporates multiple factors known to cause churn. The higher the score, the more likely the customer will churn.

CLTV: Customer Lifetime Value. A predicted CLTV is calculated using corporate formulas and existing data. The higher the value, the more valuable the customer. High value customers should be monitored for churn.

Churn Reason: A customer's specific reason for leaving the company. Directly related to Churn Category.

```
[585]: db=pd.read_excel('Telco_customer_churn.xlsx')
db
```

```
[585]:
```

	CustomerID	Count	Country	State	City	Zip Code	\
0	3668-QPYBK	1	United States	California	Los Angeles	90003	
1	9237-HQITU	1	United States	California	Los Angeles	90005	
2	9305-CDSKC	1	United States	California	Los Angeles	90006	
3	7892-P00KP	1	United States	California	Los Angeles	90010	
4	0280-XJGEX	1	United States	California	Los Angeles	90015	
...	
7038	2569-WGERO	1	United States	California	Landers	92285	
7039	6840-RESVB	1	United States	California	Adelanto	92301	
7040	2234-XADUH	1	United States	California	Amboy	92304	
7041	4801-JZAZL	1	United States	California	Angelus Oaks	92305	
7042	3186-AJIEK	1	United States	California	Apple Valley	92308	

	Lat Long	Latitude	Longitude	Gender	...	\
0	33.964131, -118.272783	33.964131	-118.272783	Male	...	
1	34.059281, -118.30742	34.059281	-118.307420	Female	...	
2	34.048013, -118.293953	34.048013	-118.293953	Female	...	
3	34.062125, -118.315709	34.062125	-118.315709	Female	...	
4	34.039224, -118.266293	34.039224	-118.266293	Male	...	
...	
7038	34.341737, -116.539416	34.341737	-116.539416	Female	...	
7039	34.667815, -117.536183	34.667815	-117.536183	Male	...	
7040	34.559882, -115.637164	34.559882	-115.637164	Female	...	
7041	34.1678, -116.86433	34.167800	-116.864330	Female	...	
7042	34.424926, -117.184503	34.424926	-117.184503	Male	...	

	Contract Paperless Billing	Payment Method	...	\
0	Month-to-month	Yes	Mailed check	
1	Month-to-month	Yes	Electronic check	
2	Month-to-month	Yes	Electronic check	
3	Month-to-month	Yes	Electronic check	
4	Month-to-month	Yes	Bank transfer (automatic)	
...	
7038	Two year	Yes	Bank transfer (automatic)	
7039	One year	Yes	Mailed check	
7040	One year	Yes	Credit card (automatic)	
7041	Month-to-month	Yes	Electronic check	
7042	Two year	Yes	Bank transfer (automatic)	

	Monthly Charges	Total Charges	Churn Label	Churn Value	Churn Score	CLTV	...	\
0	53.85	108.15	Yes	1	86	3239		
1	70.70	151.65	Yes	1	67	2701		
2	99.65	820.5	Yes	1	86	5372		
3	104.80	3046.05	Yes	1	84	5003		
4	103.70	5036.3	Yes	1	89	5340		
...		
7038	21.15	1419.4	No	0	45	5306		
7039	84.80	1990.5	No	0	59	2140		
7040	103.20	7362.9	No	0	71	5560		
7041	29.60	346.45	No	0	59	2793		
7042	105.65	6844.5	No	0	38	5097		

	Churn Reason
0	Competitor made better offer
1	Moved
2	Moved
3	Moved
4	Competitor had better devices
...	...
7038	NaN

7039	NaN
7040	NaN
7041	NaN
7042	NaN

[7043 rows x 33 columns]

Check for Null variables

[588]: db.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 33 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerID            7043 non-null   object
1   Count                  7043 non-null   int64
2   Country                7043 non-null   object
3   State                  7043 non-null   object
4   City                   7043 non-null   object
5   Zip Code               7043 non-null   int64
6   Lat Long               7043 non-null   object
7   Latitude                7043 non-null   float64
8   Longitude               7043 non-null   float64
9   Gender                 7043 non-null   object
10  Senior Citizen         7043 non-null   object
11  Partner                 7043 non-null   object
12  Dependents             7043 non-null   object
13  Tenure Months          7043 non-null   int64
14  Phone Service          7043 non-null   object
15  Multiple Lines         7043 non-null   object
16  Internet Service       7043 non-null   object
17  Online Security        7043 non-null   object
18  Online Backup          7043 non-null   object
19  Device Protection      7043 non-null   object
20  Tech Support           7043 non-null   object
21  Streaming TV           7043 non-null   object
22  Streaming Movies       7043 non-null   object
23  Contract               7043 non-null   object
24  Paperless Billing       7043 non-null   object
25  Payment Method         7043 non-null   object
26  Monthly Charges        7043 non-null   float64
27  Total Charges          7043 non-null   object
28  Churn Label            7043 non-null   object
29  Churn Value            7043 non-null   int64
30  Churn Score            7043 non-null   int64
31  CLTV                   7043 non-null   int64
```

```
32 Churn Reason      1869 non-null  object
dtypes: float64(3), int64(6), object(24)
memory usage: 1.8+ MB
```

Data Cleaning and Preprocessing

```
[591]: db.isnull().sum()
```

```
[591]: CustomerID      0
      Count          0
      Country        0
      State          0
      City           0
      Zip Code       0
      Lat Long       0
      Latitude       0
      Longitude      0
      Gender         0
      Senior Citizen 0
      Partner        0
      Dependents     0
      Tenure Months  0
      Phone Service  0
      Multiple Lines 0
      Internet Service 0
      Online Security 0
      Online Backup  0
      Device Protection 0
      Tech Support   0
      Streaming TV   0
      Streaming Movies 0
      Contract       0
      Paperless Billing 0
      Payment Method 0
      Monthly Charges 0
      Total Charges  0
      Churn Label    0
      Churn Value    0
      Churn Score    0
      CLTV           0
      Churn Reason   5174
      dtype: int64
```

So only data missing is in **Churn Reason** which tells us that 5174 kept their subscription while the rest didn't.

So, let's start with data tidying. I will check some columns in the table that in our case can be irrelevant and take them out of dataframe.

```
[595]: db['Count'].value_counts()
```

```
[595]: Count
      1      7043
      Name: count, dtype: int64
```

```
[597]: db=db.drop(columns='Count')
```

```
[599]: db['Country'].unique()
```

```
[599]: array(['United States'], dtype=object)
```

So each customer is from US which makes this column not important in our analysis so i drop it

```
[602]: db=db.drop(columns='Country')
```

```
[604]: db['State'].unique()
```

```
[604]: array(['California'], dtype=object)
```

Same thing drop it

```
[607]: db=db.drop(columns='State')
```

```
[609]: db['Senior Citizen'].unique()
```

```
[609]: array(['No', 'Yes'], dtype=object)
```

```
[611]: db['Total Charges'].unique()
print(db['Total Charges'].unique())
unique_values = db['Total Charges'].unique()
```

```
[108.15 151.65 820.5 ... 7362.9 346.45 6844.5]
```

```
[613]: # Convert 'Total Charges' to numeric, errors='coerce' will turn non-numeric
      ↪ values into NaN
      db['Total Charges'] = pd.to_numeric(db['Total Charges'], errors='coerce')

      # Handle missing values in Total Charges (NaNs)
      db['Total Charges'] = db['Total Charges'].fillna(db['Total Charges'].median())
```

Finally, we can see all possible reasons for subscription cancellation


```
[616]: db["Churn Reason"].unique()
```

```
[616]: array(['Competitor made better offer', 'Moved',
      'Competitor had better devices',
      'Competitor offered higher download speeds',
      'Competitor offered more data', 'Price too high',
      'Product dissatisfaction', 'Service dissatisfaction',
```

```
'Lack of self-service on Website', 'Network reliability',
'Limited range of services',
'Lack of affordable download/upload speed',
'Long distance charges', 'Extra data charges', "Don't know",
'Poor expertise of online support',
'Poor expertise of phone support', 'Attitude of service provider',
'Attitude of support person', 'Deceased', nan], dtype=object)
```

Now, I will create visualization of missing data, just to make point visible with which types of data I will work

```
[619]: plt.figure(figsize=(10, 6))
sns.heatmap(db.isnull(), cbar=False, cmap='viridis')
plt.title('Missing Value Visualization')
plt.show()
```



Quickly understanding the distribution and central characteristics of this data

```
[622]: db.describe()
```




```
[622]:
```

	Zip Code	Latitude	Longitude	Tenure Months	Monthly Charges \
count	7043.000000	7043.000000	7043.000000	7043.000000	7043.000000
mean	93521.964646	36.282441	-119.798880	32.371149	64.761692
std	1865.794555	2.455723	2.157889	24.559481	30.090047
min	90001.000000	32.555828	-124.301372	0.000000	18.250000
25%	92102.000000	34.030915	-121.815412	9.000000	35.500000
50%	93552.000000	36.391777	-119.730885	29.000000	70.350000
75%	95351.000000	38.224869	-118.043237	55.000000	89.850000
max	96161.000000	41.962127	-114.192901	72.000000	118.750000

	Total Charges	Churn Value	Churn Score	CLTV
count	7043.000000	7043.000000	7043.000000	7043.000000
mean	2281.916928	0.265370	58.699418	4400.295755
std	2265.270398	0.441561	21.525131	1183.057152
min	18.800000	0.000000	5.000000	2003.000000
25%	402.225000	0.000000	40.000000	3469.000000
50%	1397.475000	0.000000	61.000000	4527.000000
75%	3786.600000	1.000000	75.000000	5380.500000
max	8684.800000	1.000000	100.000000	6500.000000

I want to check scatter plot of monthly charges and total charges how does it look like

```
[625]: plt.figure(figsize=(10, 6))
sns.scatterplot(x='Monthly Charges', y='Total Charges', data=db)
plt.title('Scatter Plot of Monthly Charges vs. Total Charges')
plt.xlabel('Monthly Charges')
plt.ylabel('Total Charges')
plt.show()
```



I want to check all kinds of different distributions using histograms

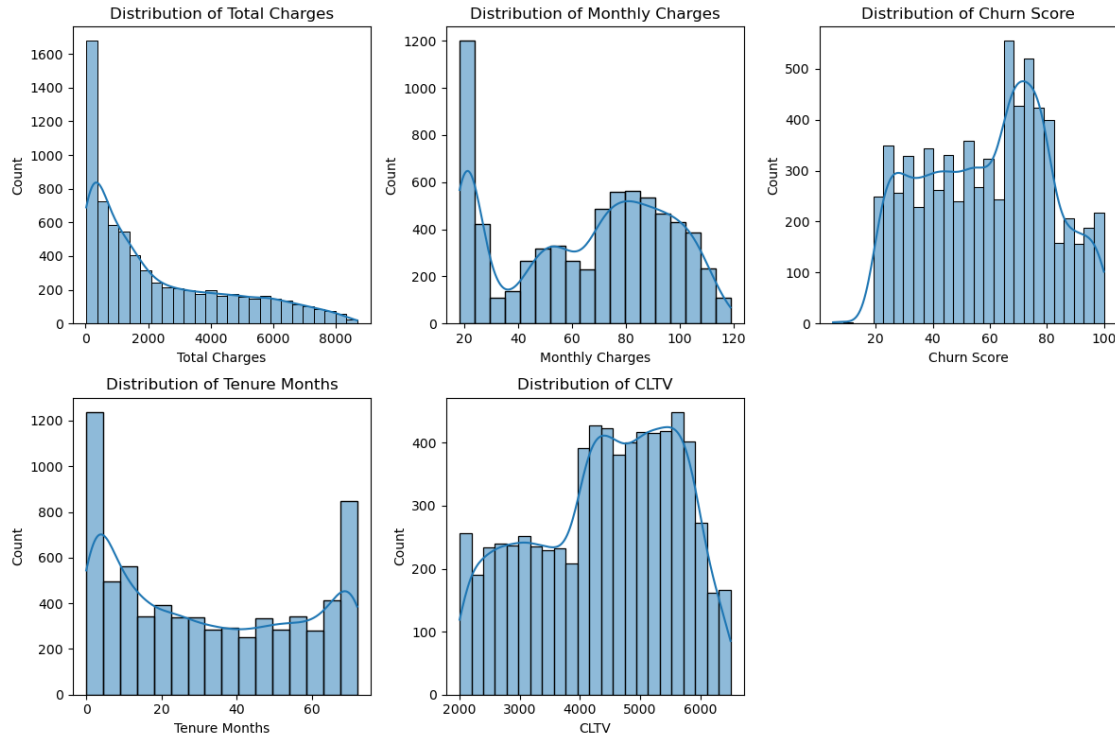
```
[628]: def distribution_numerical(data, x, index):
    plt.subplot(2, 3, index)
    sns.histplot(data=data, x=x, kde=True)
    plt.title(f"Distribution of {x}")
    plt.tight_layout()

    numerical_columns = ['Total Charges', 'Monthly Charges', 'Churn Score', 'Tenure_
↳Months', 'CLTV']

    fig = plt.figure(figsize=(12, 8))

    for index, col in enumerate(numerical_columns):
        distribution_numerical(db, col, index + 1)

    plt.show()
```



Now i want to plot most popular reason of cancelling subscription:

```
[630]: churn_reason_counts = db['Churn Reason'].dropna().value_counts()

plt.figure(figsize=(10, 6))
bars = plt.bar(churn_reason_counts.index, churn_reason_counts.values,
               color='skyblue')
plt.xlabel("Churn Reason")
plt.ylabel("Count")
plt.title("Most Popular Churn Reasons")
plt.xticks(rotation=90)

total_count = churn_reason_counts.sum()

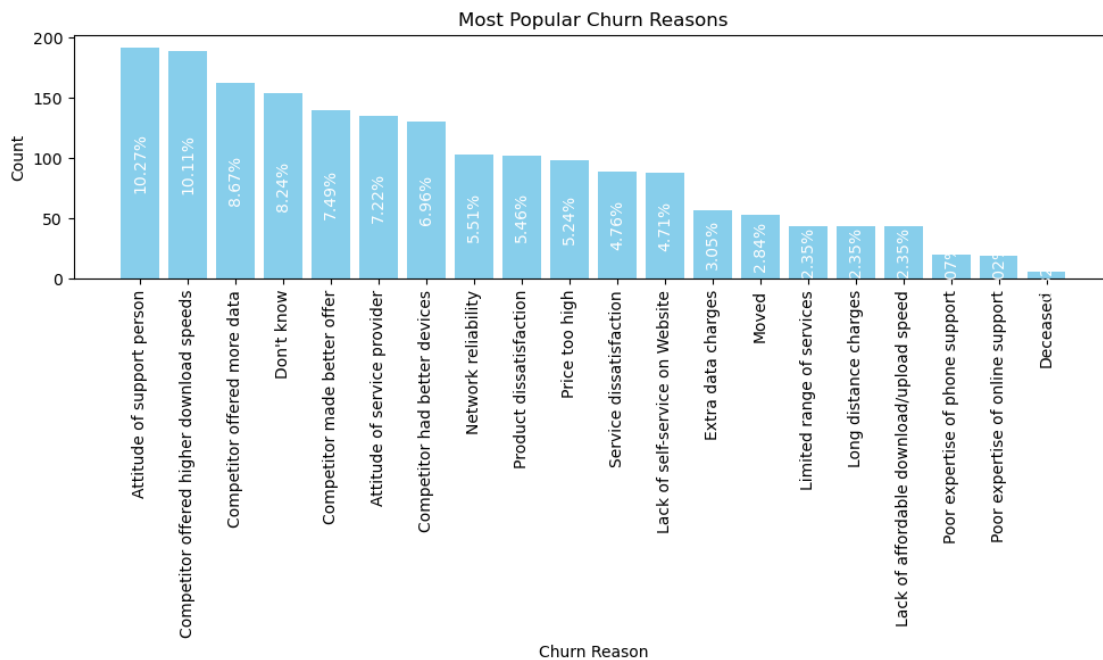
for bar, count in zip(bars, churn_reason_counts.values):
    height = bar.get_height()
    percentage = (count / total_count) * 100
    plt.text(
        bar.get_x() + bar.get_width() / 2,
        height / 2,
```

```

        f'{percentage:.2f}%',
        ha='center',
        va='center',
        color='white',
        fontsize=10,
        rotation=90
    )

plt.tight_layout()
plt.show()

```



We can clearly see why we losing customers Let's find out about which amounts in losses it could lead

```

[634]: #averages
avg_total_charges = db['Total Charges'].mean()
avg_monthly_charges = db['Monthly Charges'].mean()

print("Average Total Charges:", avg_total_charges, "$")
print("Average Monthly Charges:", avg_monthly_charges, "$")

null_percentage = (db["Churn Reason"].isnull().sum() / len(db)) * 100
cancel=(len(db)-db["Churn Reason"].isnull().sum())
print("Average loss on Monthly Charges:",cancel*avg_monthly_charges, "$")
print("Average loss on Total Charges:",cancel*avg_total_charges, "$")

```

Average Total Charges: 2281.9169281556156 \$
Average Monthly Charges: 64.76169246059918 \$
Average loss on Monthly Charges: 121039.60320885986 \$
Average loss on Total Charges: 4264902.738722846 \$

We loosing such amount of money because of Attitude of support person as a reason number 1! then we have: competitors,download speeds,less data,attitude of service provider,product dissatisfaction,service dissatisfaction,price etc.

Generating Map with spread of each customer to see which types of contract are most likely to become a churn value using plotly

```
[638]: def generate_map_trace(data, color, width, height, title):

    fig = px.scatter_mapbox(
        data,
        lat='Latitude',
        lon='Longitude',
        color=color,
        hover_name='CustomerID',
        zoom=4,
        height=height,
        width=width,
        title=title
    )

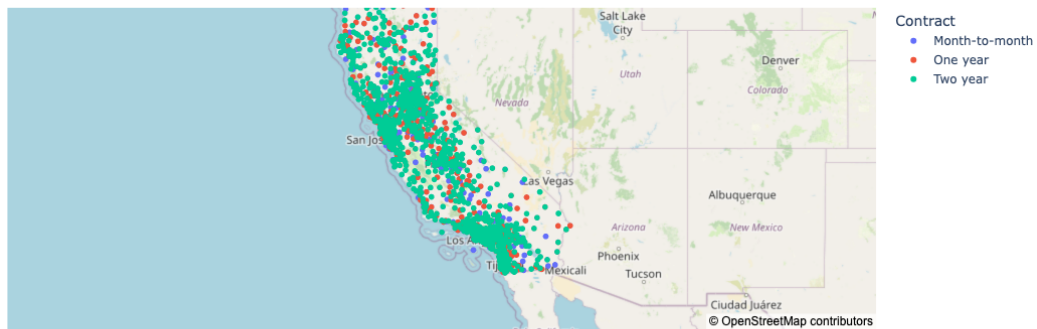
    fig.update_layout(mapbox_style="open-street-map")
    fig.show()

churn_value = pd.DataFrame()
churn_value["Churn Value"] = db["Churn Value"].map({0: "No Churn", 1: "Churn"})
churn_value[[col for col in db.columns if col != 'Churn Value']] = db[[col for
    ↪col in db.columns if col != 'Churn Value']]


color = 'Contract'

generate_map_trace(
    churn_value[churn_value['Churn Value'] == 'No Churn'],
    color=color, width=900, height=500, title='Map with No Churn Values'
)
generate_map_trace(
    churn_value[churn_value['Churn Value'] == 'Churn'],
    color=color, width=900, height=500, title='Map with Only Churn Values'
)
```

Map with No Churn Values



Map with Only Churn Values



From this map is clearly visible that Customers with two year and one year contracts almost never canceling and provide company with most amount of money which is obvious.

```
[641]: # Counting customers who complete the purchase versus those who do not
churn_counts = db['Churn Label'].value_counts()
print("Count of customers who complete the purchase vs those who churn:")
print(churn_counts)

# Calculate average spending for customers based on churn status
average_charges_by_churn = db.groupby('Churn Label')['Total Charges'].mean()
print("\nAverage spending by churn status:")
print(average_charges_by_churn)
```

```

# Calculate average spending based on tenure for loyal and non-loyal customers
↳ to see dependence
average_charges_by_tenure = db.groupby(['Churn Label', 'Tenure Months'])['Total_
↳ Charges'].mean().reset_index()
print("\nAverage spending by tenure for churned and non-churned customers:")
print(average_charges_by_tenure)

#Creating line plot based on that info
db = db.replace([np.inf, -np.inf], np.nan)
plt.figure(figsize=(12, 8))
sns.lineplot(data=average_charges_by_tenure, x='Tenure Months', y='Total_
↳ Charges', hue='Churn Label')
plt.title('Average Spending by Tenure for Churned and Non-Churned Customers')
plt.xlabel('Tenure')
plt.ylabel('Average Spending (Total Charges)')
plt.legend(title='Churn Status')
plt.show()

```

Count of customers who complete the purchase vs those who churn:

Churn Label

No 5174

Yes 1869

Name: count, dtype: int64

Average spending by churn status:

Churn Label

No 2552.882494

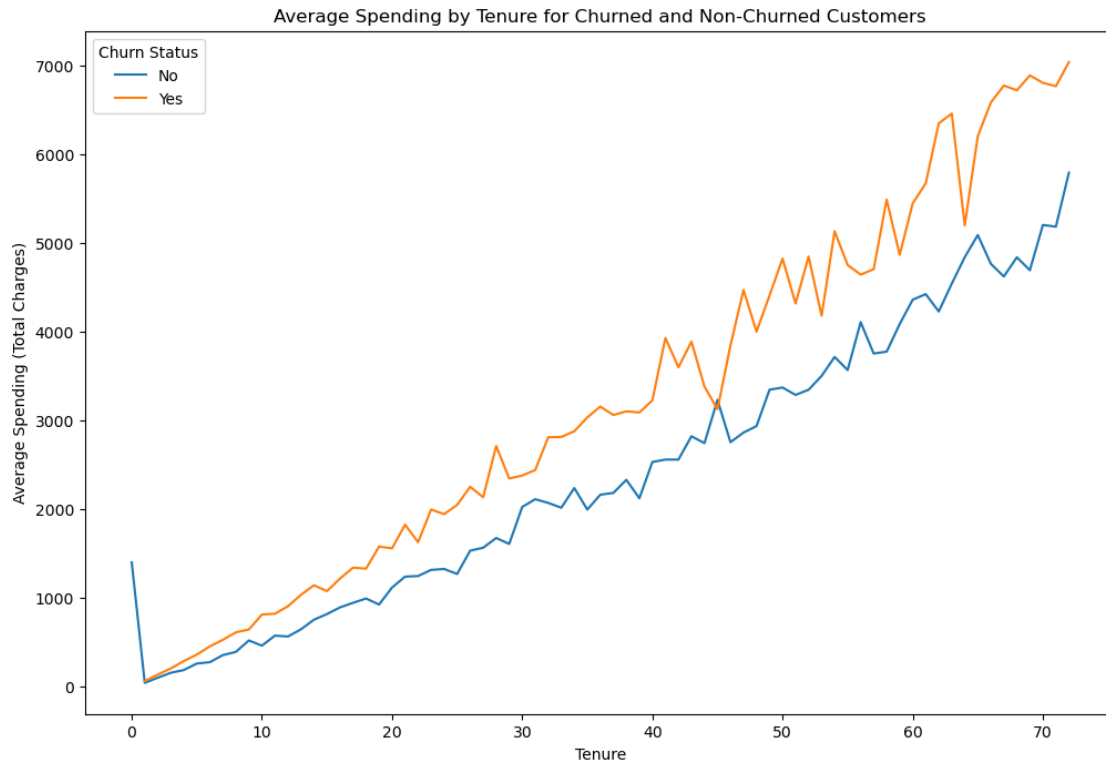
Yes 1531.796094

Name: Total Charges, dtype: float64

Average spending by tenure for churned and non-churned customers:

	Churn Label	Tenure Months	Total Charges
0	No	0	1397.475000
1	No	1	37.909013
2	No	2	95.997391
3	No	3	152.135849
4	No	4	182.525806
..
140	Yes	68	6720.550000
141	Yes	69	6887.931250
142	Yes	70	6803.995455
143	Yes	71	6765.908333
144	Yes	72	7039.150000

[145 rows x 3 columns]



We can see that loyal customers paying almost same amount as churned ones! It looks not that obvious, but to set up the contract they pay I assume some kind of fee which leads in huge spending on first month then it comesback to normal.

Finding out more about churned customers

```
[645]: # Filter for churned customers
churned_customers = db[db['Churn Label'] == 'Yes']

# average tenure in months for churned customers
avg_tenure = churned_customers['Tenure Months'].mean()
print("Average Tenure Months for Churned Customers:", avg_tenure)

# finding most popular tenure month among churned customers
most_popular_tenure = churned_customers['Tenure Months'].value_counts().idxmax()
popular_tenure_count = churned_customers['Tenure Months'].value_counts().max()
print("Most Popular Tenure Month for Churned Customers:", most_popular_tenure,
      f"({popular_tenure_count} occurrences)")
```

Average Tenure Months for Churned Customers: 17.979133226324237

Most Popular Tenure Month for Churned Customers: 1 (380 occurrences)

That's very important it tells us that most customers leave in the first month of their month-to-

month subscription because of the reasons above. Now, I can find out the correlation on this matter between the reasons, which will allow the company to fix it which potentially could lead to great success because as we can see even for churned customers, the average tenure months is almost 18!!! So that means if we can hold our new customers to more than 1 month that could lead to a higher probability of prolonging their stay up to 17 months and more! Which on average could lead to 2175081.6\$ in profit!

```
[648]: # Correlation Matrix
numeric_cols = db.select_dtypes(include=[np.number])
corr_matrix = numeric_cols.corr()
print("Correlation Matrix:")
print(corr_matrix)

# Plot a heatmap of correlation matrix
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title("Correlation Matrix")
plt.show()
```

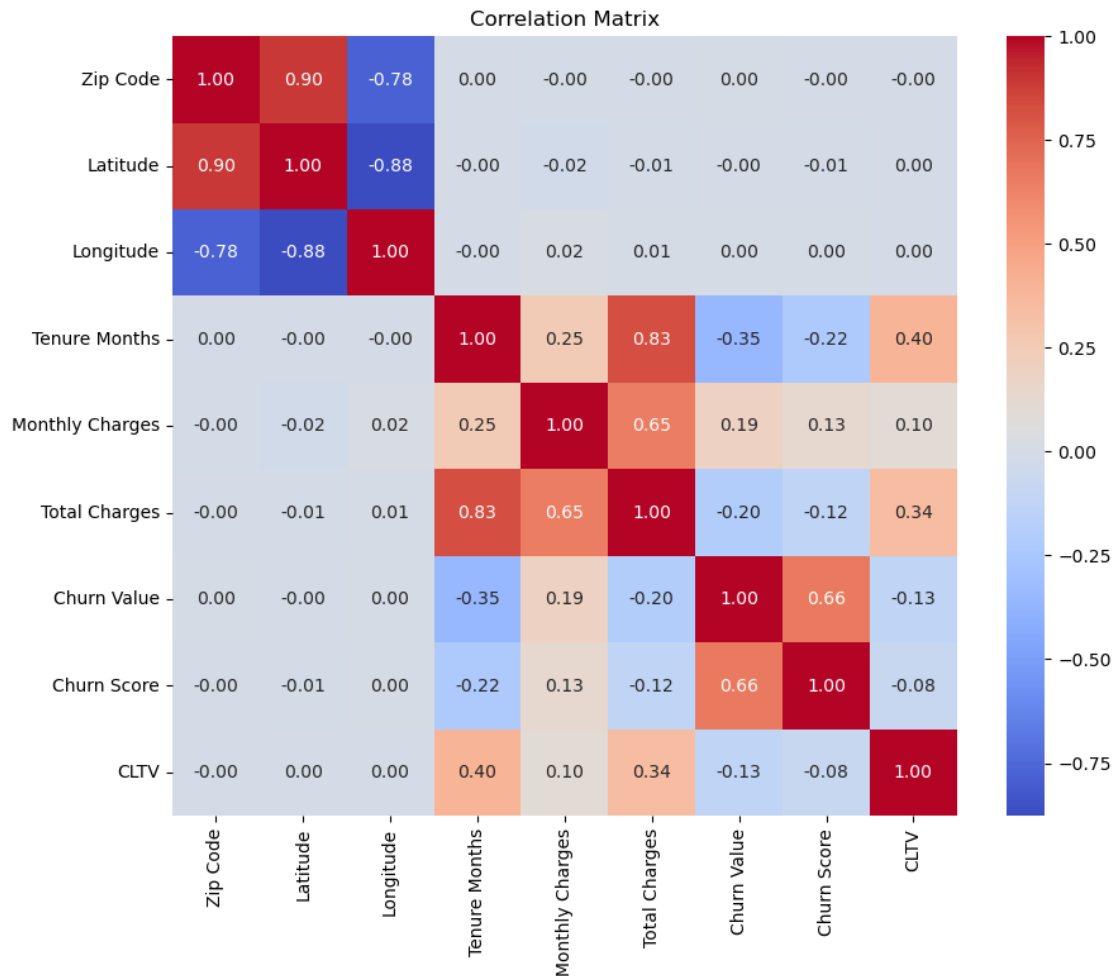
Correlation Matrix:

	Zip Code	Latitude	Longitude	Tenure Months	\
Zip Code	1.000000	0.895743	-0.784271	0.001041	
Latitude	0.895743	1.000000	-0.876779	-0.001631	
Longitude	-0.784271	-0.876779	1.000000	-0.001678	
Tenure Months	0.001041	-0.001631	-0.001678	1.000000	
Monthly Charges	-0.004596	-0.019899	0.024098	0.247900	
Total Charges	-0.001237	-0.010168	0.008977	0.825464	
Churn Value	0.003346	-0.003384	0.004594	-0.352229	
Churn Score	-0.002769	-0.007684	0.004260	-0.224987	
CLTV	-0.003562	0.000886	0.000485	0.396406	

	Monthly Charges	Total Charges	Churn Value	Churn Score	\
Zip Code	-0.004596	-0.001237	0.003346	-0.002769	
Latitude	-0.019899	-0.010168	-0.003384	-0.007684	
Longitude	0.024098	0.008977	0.004594	0.004260	
Tenure Months	0.247900	0.825464	-0.352229	-0.224987	
Monthly Charges	1.000000	0.650864	0.193356	0.133754	
Total Charges	0.650864	1.000000	-0.199037	-0.123948	
Churn Value	0.193356	-0.199037	1.000000	0.664897	
Churn Score	0.133754	-0.123948	0.664897	1.000000	
CLTV	0.098693	0.341723	-0.127463	-0.079782	

	CLTV
Zip Code	-0.003562
Latitude	0.000886
Longitude	0.000485
Tenure Months	0.396406
Monthly Charges	0.098693

Total Charges 0.341723
 Churn Value -0.127463
 Churn Score -0.079782
 CLTV 1.000000



From that we can see that Tenure months have a strong correlation between total charges and CLTV Monthly charges have strong correlation with Total charges Total charges have strong correlation between Tenure Months,Monthly charges and CLTV CLTV have medium correlation between Tenure Months and Total charges Churn Value and Churn Score are opposites of each other

```
[651]: # Chi-Square Test: Churn Reason vs Churn Label
# contingency table
contingency_table = pd.crosstab(db['Churn Reason'], db['Churn Label'])
print("\nContingency Table (Churn Reason vs. Churn Label):")

# Extracting count for yes
```

```

counts = contingency_table['Yes']
total_count = counts.sum()

# df with churn reason, count, and percentage
df_counts = counts.reset_index()
df_counts.columns = ['Churn Reason', 'Count']
df_counts['Percentage'] = (df_counts['Count'] / total_count) * 100

df_counts = df_counts.sort_values(by='Count', ascending=False)

print("\nChurn Reasons with Count and Percentage:")
print(df_counts)

# chi-square test
chi2, p, dof, expected = stats.chi2_contingency(contingency_table)
print(f"\nChi-square Test (Churn Reason vs. Churn Label):")
print(f"  Chi-square Statistic: {chi2:.2f}")
print(f"  p-value: {p:.4f}")
print(f"  Degrees of Freedom: {dof}")

```

Contingency Table (Churn Reason vs. Churn Label):

Churn Reasons with Count and Percentage:

	Churn Reason	Count	Percentage
1	Attitude of support person	192	10.272873
4	Competitor offered higher download speeds	189	10.112360
5	Competitor offered more data	162	8.667737
7	Don't know	154	8.239700
3	Competitor made better offer	140	7.490637
0	Attitude of service provider	135	7.223114
2	Competitor had better devices	130	6.955591
14	Network reliability	103	5.510968
18	Product dissatisfaction	102	5.457464
17	Price too high	98	5.243446
19	Service dissatisfaction	89	4.761905
10	Lack of self-service on Website	88	4.708400
8	Extra data charges	57	3.049759
13	Moved	53	2.835741
11	Limited range of services	44	2.354200
12	Long distance charges	44	2.354200
9	Lack of affordable download/upload speed	44	2.354200
16	Poor expertise of phone support	20	1.070091
15	Poor expertise of online support	19	1.016586
6	Deceased	6	0.321027

Chi-square Test (Churn Reason vs. Churn Label):

Chi-square Statistic: 0.00
p-value: 1.0000
Degrees of Freedom: 0

High p-value: Suggests that any differences are likely due to random variation, and we do not have enough evidence to conclude that an association exists.

```
[654]: #the Most Popular Tenure among Churned Customers

churned_customers = db[db['Churn Label'] == 'Yes']
most_popular_tenure = churned_customers['Tenure Months'].value_counts().idxmax()
print(f"\nMost Popular Tenure Months among churned customers:␣
↳{most_popular_tenure}")

db['Popular Tenure'] = (db['Tenure Months'] == most_popular_tenure)


# contingency table
contingency_table2 = pd.crosstab(db['Churn Reason'], db['Popular Tenure'])
print("\nContingency Table (Churn Reason vs. Popular Tenure):")
popular_tenure_df = db[db['Popular Tenure'] == True]


# count of the occurrences
churn_reason_counts = popular_tenure_df['Churn Reason'].value_counts().
↳reset_index()
churn_reason_counts.columns = ['Churn Reason', 'Count']
churned_customers = churned_customers.copy()
churned_customers['Popular Tenure'] = (churned_customers['Tenure Months'] ==␣
↳most_popular_tenure)

print("Total churned customers with Popular Tenure True:
↳",churned_customers['Popular Tenure'].sum())

churn_reason_counts['Percentage'] = (churn_reason_counts['Count'] /␣
↳churned_customers['Popular Tenure'].sum()) * 100

churn_reason_counts = churn_reason_counts.sort_values(by='Count',␣
↳ascending=False)
print(churn_reason_counts)


#Chi-Square Test: Churn Reason vs. Popular Tenure
```

```
chi2_2, p_2, dof_2, expected_2 = stats.chi2_contingency(contingency_table2)
print(f"\nChi-square Test (Churn Reason vs. Popular Tenure):")
print(f"  Chi-square Statistic: {chi2_2:.2f}")
print(f"  p-value: {p_2:.4f}")
print(f"  Degrees of Freedom: {dof_2}")
```

Most Popular Tenure Months among churned customers: 1

Contingency Table (Churn Reason vs. Popular Tenure):

Total churned customers with Popular Tenure True: 380

	Churn Reason	Count	Percentage
0	Attitude of support person	53	13.947368
1	Don't know	33	8.684211
2	Competitor made better offer	31	8.157895
3	Attitude of service provider	30	7.894737
4	Competitor offered higher download speeds	28	7.368421
5	Competitor had better devices	28	7.368421
6	Competitor offered more data	22	5.789474
9	Service dissatisfaction	21	5.526316
10	Price too high	21	5.526316
8	Network reliability	21	5.526316
7	Product dissatisfaction	21	5.526316
11	Lack of self-service on Website	17	4.473684
12	Moved	10	2.631579
13	Limited range of services	9	2.368421
14	Extra data charges	8	2.105263
15	Lack of affordable download/upload speed	8	2.105263
16	Poor expertise of phone support	6	1.578947
17	Long distance charges	6	1.578947
18	Poor expertise of online support	5	1.315789
19	Deceased	2	0.526316

Chi-square Test (Churn Reason vs. Popular Tenure):

Chi-square Statistic: 20.92

p-value: 0.3413

Degrees of Freedom: 19

A Chi-square Statistic of 20.92 is high which would suggest a large discrepancy. A p-value of 0.3413 is well above the common significance threshold of 0.05. This means that the observed differences are likely due to random, and there isn't strong evidence to reject the null hypothesis.

From this Contingency Table we can clearly see that we can avoid some reasons of cancelation by training company employees,offer higher download speeds,offer more data,update current devices,add self services on website etc.Which will lead for longer subscription periods of customers!

So as we know that not fixing this issues lead to a losses of company monthly almost 25k\$ for customers who got dissatisfied at first month,and 125k\$ on average monthly!!!

```
[659]: print("Average loss on new customers:", churned_customers['Popular Tenure'].
        ↪sum()*avg_monthly_charges, "$")
```

Average loss on new customers: 24609.443135027686 \$

```
[661]: print("Average loss on Monthly Charges:", cancel*avg_monthly_charges, "$")
```

Average loss on Monthly Charges: 121039.60320885986 \$

Let's find out is there any proves that amount of monthly charges can affect the duration of subscription:

```
[664]: avg_monthly = db.groupby("Contract")["Monthly Charges"].mean()
print("Average Monthly Charges by Contract Type:")
print(avg_monthly)

#monthly charges for each contract type
month_to_month = db[db["Contract"] == "Month-to-month"]["Monthly Charges"]
one_year = db[db["Contract"] == "One year"]["Monthly Charges"]
two_year = db[db["Contract"] == "Two year"]["Monthly Charges"]

#one-way ANOVA test to compare the three groups
f_stat, p_val = stats.f_oneway(month_to_month, one_year, two_year)
print("\nANOVA Test Results (Monthly Charges by Contract):")
print("F-statistic:", f_stat)
print("p-value:", p_val)

#checking if monthly Charges affect the duration of subscription

#Pearson correlation between Monthly Charges and Tenure Months
corr_coef = db['Monthly Charges'].corr(db['Tenure Months'])
print("\nPearson Correlation between Monthly Charges and Tenure Months:",
        ↪corr_coef)

plt.figure(figsize=(10,6))
sns.regplot(x='Monthly Charges', y='Tenure Months', data=db,
            ↪scatter_kws={'alpha':0.5})
plt.title('Relationship between Monthly Charges and Tenure Months')
plt.xlabel('Monthly Charges')
plt.ylabel('Tenure Months')
plt.tight_layout()
plt.show()

#simple linear regression using OLS
X = db['Monthly Charges']
y = db['Tenure Months']
X = sm.add_constant(X) # Adds an intercept term to the model
```

```

model = sm.OLS(y, X).fit()
print("\nLinear Regression Results:")
print(model.summary())

# Since Monthly Charges starts around $55, re-center the predictor
db['Monthly Charges Centered'] = db['Monthly Charges'] - 55

# Fit the model using the centered predictor
X_centered = db['Monthly Charges Centered']
X_centered = sm.add_constant(X_centered)
model_centered = sm.OLS(y, X_centered).fit()

print("\nLinear Regression Results (Centered at $55):")
print(model_centered.summary())

```

Average Monthly Charges by Contract Type:

Contract

Month-to-month 66.398490

One year 65.048608

Two year 60.770413

Name: Monthly Charges, dtype: float64

ANOVA Test Results (Monthly Charges by Contract):

F-statistic: 20.828045474730278

p-value: 9.575270975935035e-10

Pearson Correlation between Monthly Charges and Tenure Months:

0.24789985628615008



Linear Regression Results:

OLS Regression Results

=====					
Dep. Variable:	Tenure Months	R-squared:	0.061		
Model:	OLS	Adj. R-squared:	0.061		
Method:	Least Squares	F-statistic:	461.0		
Date:	Sat, 05 Apr 2025	Prob (F-statistic):	4.09e-99		
Time:	11:21:33	Log-Likelihood:	-32315.		
No. Observations:	7043	AIC:	6.463e+04		
Df Residuals:	7041	BIC:	6.465e+04		
Df Model:	1				
Covariance Type:	nonrobust				
=====					
===					
	coef	std err	t	P> t	[0.025
0.975]					

const	19.2675	0.673	28.633	0.000	17.948
20.587					
Monthly Charges	0.2023	0.009	21.472	0.000	0.184
0.221					
=====					
Omnibus:	7646.995	Durbin-Watson:	1.662		
Prob(Omnibus):	0.000	Jarque-Bera (JB):	513.504		

Skew:	0.252	Prob(JB):	3.12e-112
Kurtosis:	1.777	Cond. No.	170.

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Linear Regression Results (Centered at \$55):

OLS Regression Results

=====

Dep. Variable:	Tenure Months	R-squared:	0.061
Model:	OLS	Adj. R-squared:	0.061
Method:	Least Squares	F-statistic:	461.0
Date:	Sat, 05 Apr 2025	Prob (F-statistic):	4.09e-99
Time:	11:21:33	Log-Likelihood:	-32315.
No. Observations:	7043	AIC:	6.463e+04
Df Residuals:	7041	BIC:	6.465e+04
Df Model:	1		
Covariance Type:	nonrobust		

=====

=====

	coef	std err	t	P> t	[0.025
0.975]					

const	30.3960	0.298	101.973	0.000	29.812
30.980					
Monthly Charges Centered	0.2023	0.009	21.472	0.000	0.184
0.221					

=====

Omnibus:	7646.995	Durbin-Watson:	1.662
Prob(Omnibus):	0.000	Jarque-Bera (JB):	513.504
Skew:	0.252	Prob(JB):	3.12e-112
Kurtosis:	1.777	Cond. No.	33.3

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

ANNOVA test which compares the means of the three contract types to see if at least one group has a significantly different mean monthly charge The F-statistic of about 20.83 indicates a significant difference in the variability between groups compared to the variability within groups. The p-value is extremely small (around 9.6e-10), which is much less than the common **significance level** 0.05 These results provide strong statistical evidence that the average monthly charges differ by contract type.

A **correlation coefficient** of approximately 0.25 indicates a positive but modest linear relationship

between monthly charges and the duration of the subscription. This suggests that there is a slight tendency for customers who pay lower monthly charges to have longer tenures, but the relationship isn't very strong. Many other factors likely influence how long customers stay.

OLS Model:

const = 19.27: If a customer were charged 55\$ per month a hypothetical scenario, the model predicts they would have a tenure of about 30.39 months. It serves as the baseline for the model. Coefficient for Monthly Charges 0.2023: For every \$1 increase in monthly charges, the model predicts an increase of about 0.2023 months in tenure. (Which is slightly suspicious)

R-squared 0.061: This tells us that approximately 6.1% of the variance in Tenure Months is explained by Monthly Charges.

Overall: At \$55 per month, the predicted customer tenure is about 30.4 months. For every extra dollar above \$55, customer tenure increases by about 0.20 months, on average.

The F-statistic and its associated very low p-value < 0.001 indicates that this relationship is statistically significant. Although the effect is statistically significant, monthly charges alone explain only a small portion of the variability in subscription duration. Other factors likely have a larger impact on tenure.

2.2.1 Building Predictive model

```
[668]: db.head()
```

```
[668]:
```

	CustomerID	City	Zip Code	Lat	Long	Latitude	\
0	3668-QPYBK	Los Angeles	90003	33.964131,	-118.272783	33.964131	
1	9237-HQITU	Los Angeles	90005	34.059281,	-118.30742	34.059281	
2	9305-CDSKC	Los Angeles	90006	34.048013,	-118.293953	34.048013	
3	7892-P00KP	Los Angeles	90010	34.062125,	-118.315709	34.062125	
4	0280-XJGEX	Los Angeles	90015	34.039224,	-118.266293	34.039224	

	Longitude	Gender	Senior Citizen	Partner	Dependents	...	\
0	-118.272783	Male	No	No	No	...	
1	-118.307420	Female	No	No	Yes	...	
2	-118.293953	Female	No	No	Yes	...	
3	-118.315709	Female	No	Yes	Yes	...	
4	-118.266293	Male	No	No	Yes	...	

	Payment Method	Monthly Charges	Total Charges	Churn Label	\
0	Mailed check	53.85	108.15	Yes	
1	Electronic check	70.70	151.65	Yes	
2	Electronic check	99.65	820.50	Yes	
3	Electronic check	104.80	3046.05	Yes	
4	Bank transfer (automatic)	103.70	5036.30	Yes	

	Churn Value	Churn Score	CLTV	Churn Reason	Popular	Tenure	\
0	1	86	3239	Competitor made better offer		False	
1	1	67	2701	Moved		False	

2	1	86	5372		Moved	False
3	1	84	5003		Moved	False
4	1	89	5340	Competitor had better devices		False

Monthly Charges Centered	
0	-1.15
1	15.70
2	44.65
3	49.80
4	48.70

[5 rows x 32 columns]

```
[670]: from sklearn.preprocessing import LabelEncoder, StandardScaler

categorical_columns = [
    'Gender', 'Senior Citizen', 'Partner', 'Dependents', 'Phone Service',
    'Multiple Lines', 'Internet Service', 'Online Security', 'Online Backup',
    'Device Protection', 'Tech Support', 'Streaming TV', 'Streaming Movies',
    'Contract', 'Paperless Billing', 'Payment Method', 'Churn Reason'
]

label_encoder = LabelEncoder()

for column in categorical_columns:
    db[column] = label_encoder.fit_transform(db[column])

print(db.head())
```

	CustomerID	City	Zip Code	Lat	Long	Latitude \
0	3668-QPYBK	Los Angeles	90003	33.964131,	-118.272783	33.964131
1	9237-HQITU	Los Angeles	90005	34.059281,	-118.30742	34.059281
2	9305-CDSKC	Los Angeles	90006	34.048013,	-118.293953	34.048013
3	7892-POOKP	Los Angeles	90010	34.062125,	-118.315709	34.062125
4	0280-XJGEX	Los Angeles	90015	34.039224,	-118.266293	34.039224

	Longitude	Gender	Senior Citizen	Partner	Dependents	...	\
0	-118.272783	1	0	0	0	...	
1	-118.307420	0	0	0	1	...	
2	-118.293953	0	0	0	1	...	
3	-118.315709	0	0	1	1	...	
4	-118.266293	1	0	0	1	...	

	Payment Method	Monthly Charges	Total Charges	Churn Label	Churn Value \
0	3	53.85	108.15	Yes	1
1	2	70.70	151.65	Yes	1
2	2	99.65	820.50	Yes	1

3	2	104.80	3046.05	Yes	1
4	0	103.70	5036.30	Yes	1

	Churn Score	CLTV	Churn Reason	Popular	Tenure	Monthly Charges Centered
0	86	3239	3		False	-1.15
1	67	2701	13		False	15.70
2	86	5372	13		False	44.65
3	84	5003	13		False	49.80
4	89	5340	2		False	48.70

[5 rows x 32 columns]

```
[672]: X = db.drop(columns=['Churn Label', 'CustomerID', 'Churn Value','CustomerID',
↪ 'City', 'Zip Code', 'Lat Long', 'Churn Reason', 'Monthly Charges Centered'])

y = db['Churn Label']
```

[674]: X

```
[674]:      Latitude  Longitude  Gender  Senior Citizen  Partner  Dependents  \
0    33.964131 -118.272783      1              0        0          0
1    34.059281 -118.307420      0              0        0          1
2    34.048013 -118.293953      0              0        0          1
3    34.062125 -118.315709      0              0        1          1
4    34.039224 -118.266293      1              0        0          1
...      ...      ...      ...      ...      ...      ...
7038  34.341737 -116.539416      0              0        0          0
7039  34.667815 -117.536183      1              0        1          1
7040  34.559882 -115.637164      0              0        1          1
7041  34.167800 -116.864330      0              0        1          1
7042  34.424926 -117.184503      1              0        0          0
```

	Tenure Months	Phone Service	Multiple Lines	Internet Service	...	\
0	2	1	0	0	...	
1	2	1	0	1	...	
2	8	1	2	1	...	
3	28	1	2	1	...	
4	49	1	2	1	...	
...	
7038	72	1	0	2	...	
7039	24	1	2	0	...	
7040	72	1	2	1	...	
7041	11	0	1	0	...	
7042	66	1	0	1	...	

	Streaming TV	Streaming Movies	Contract	Paperless Billing	\
0	0	0	0	1	

1	0	0	0	1
2	2	2	0	1
3	2	2	0	1
4	2	2	0	1
...
7038	1	1	2	1
7039	2	2	1	1
7040	2	2	1	1
7041	0	0	0	1
7042	2	2	2	1

	Payment Method	Monthly Charges	Total Charges	Churn Score	CLTV \
0	3	53.85	108.15	86	3239
1	2	70.70	151.65	67	2701
2	2	99.65	820.50	86	5372
3	2	104.80	3046.05	84	5003
4	0	103.70	5036.30	89	5340
...
7038	0	21.15	1419.40	45	5306
7039	3	84.80	1990.50	59	2140
7040	1	103.20	7362.90	71	5560
7041	2	29.60	346.45	59	2793
7042	0	105.65	6844.50	38	5097

	Popular Tenure
0	False
1	False
2	False
3	False
4	False
...	...
7038	False
7039	False
7040	False
7041	False
7042	False

[7043 rows x 24 columns]

```
[676]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
↳ random_state=42)
X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 24 columns):
#   Column                Non-Null Count  Dtype
---

```

```

0   Latitude          7043 non-null   float64
1   Longitude         7043 non-null   float64
2   Gender            7043 non-null   int64
3   Senior Citizen    7043 non-null   int64
4   Partner           7043 non-null   int64
5   Dependents        7043 non-null   int64
6   Tenure Months     7043 non-null   int64
7   Phone Service     7043 non-null   int64
8   Multiple Lines    7043 non-null   int64
9   Internet Service  7043 non-null   int64
10  Online Security   7043 non-null   int64
11  Online Backup     7043 non-null   int64
12  Device Protection 7043 non-null   int64
13  Tech Support      7043 non-null   int64
14  Streaming TV      7043 non-null   int64
15  Streaming Movies  7043 non-null   int64
16  Contract          7043 non-null   int64
17  Paperless Billing  7043 non-null   int64
18  Payment Method    7043 non-null   int64
19  Monthly Charges   7043 non-null   float64
20  Total Charges     7043 non-null   float64
21  Churn Score       7043 non-null   int64
22  CLTV              7043 non-null   int64
23  Popular Tenure    7043 non-null   bool
dtypes: bool(1), float64(4), int64(19)
memory usage: 1.2 MB

```

[678]: `y.info()`

```

<class 'pandas.core.series.Series'>
RangeIndex: 7043 entries, 0 to 7042
Series name: Churn Label
Non-Null Count  Dtype
-----
7043 non-null   object
dtypes: object(1)
memory usage: 55.2+ KB

```

```

[680]: from sklearn.metrics import accuracy_score, confusion_matrix, ConfusionMatrixDisplay
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB

models = {
    "Decision Tree": DecisionTreeClassifier(),

```

```

    "Naive Bayes": GaussianNB(),
    "Random Forest": RandomForestClassifier(),
}

results = {}
for name, model in models.items():
    model.fit(X_train, y_train)

    y_pred_train = model.predict(X_train)
    y_pred_test = model.predict(X_test)

    train_accuracy = accuracy_score(y_train, y_pred_train)
    test_accuracy = accuracy_score(y_test, y_pred_test)


    confusion = confusion_matrix(y_test, y_pred_test)


    results[name] = {
        "train_accuracy": train_accuracy,
        "test_accuracy": test_accuracy,
        "confusion_matrix": confusion
    }


    disp = ConfusionMatrixDisplay(confusion_matrix=confusion)
    disp.plot(cmap=plt.cm.Blues)
    plt.title(f'Confusion Matrix for {name}')
    plt.show()

for name, metrics in results.items():
    print(f"{name}:")
    print(f"  Train Accuracy: {metrics['train_accuracy']:.4f}")
    print(f"  Test Accuracy: {metrics['test_accuracy']:.4f}")
    print("-" * 40)

```







Decision Tree:

Train Accuracy: 1.0000

Test Accuracy: 0.9049

Naive Bayes:

Train Accuracy: 0.8803

Test Accuracy: 0.8680

Random Forest:

Train Accuracy: 1.0000

Test Accuracy: 0.9309

```
[537]: y_churn_reason = db['Churn Reason']
X = db.drop(columns=['Churn Label', 'Churn Reason', 'CustomerID', 'Churn_
↳Value', 'CustomerID', 'City', 'Zip Code', 'Lat Long', 'Monthly Charges_
↳Centered'])
```

```

X_train, X_test, y_train, y_test = train_test_split(X, y_churn_reason,
    ↪test_size=0.3, random_state=42)

models = {
    "Decision Tree": DecisionTreeClassifier(),
    "Naive Bayes": GaussianNB(),
    "Random Forest": RandomForestClassifier(),
}

results = {}
for name, model in models.items():
    model.fit(X_train, y_train)

    y_pred_train = model.predict(X_train)
    y_pred_test = model.predict(X_test)

    train_accuracy = accuracy_score(y_train, y_pred_train)
    test_accuracy = accuracy_score(y_test, y_pred_test)

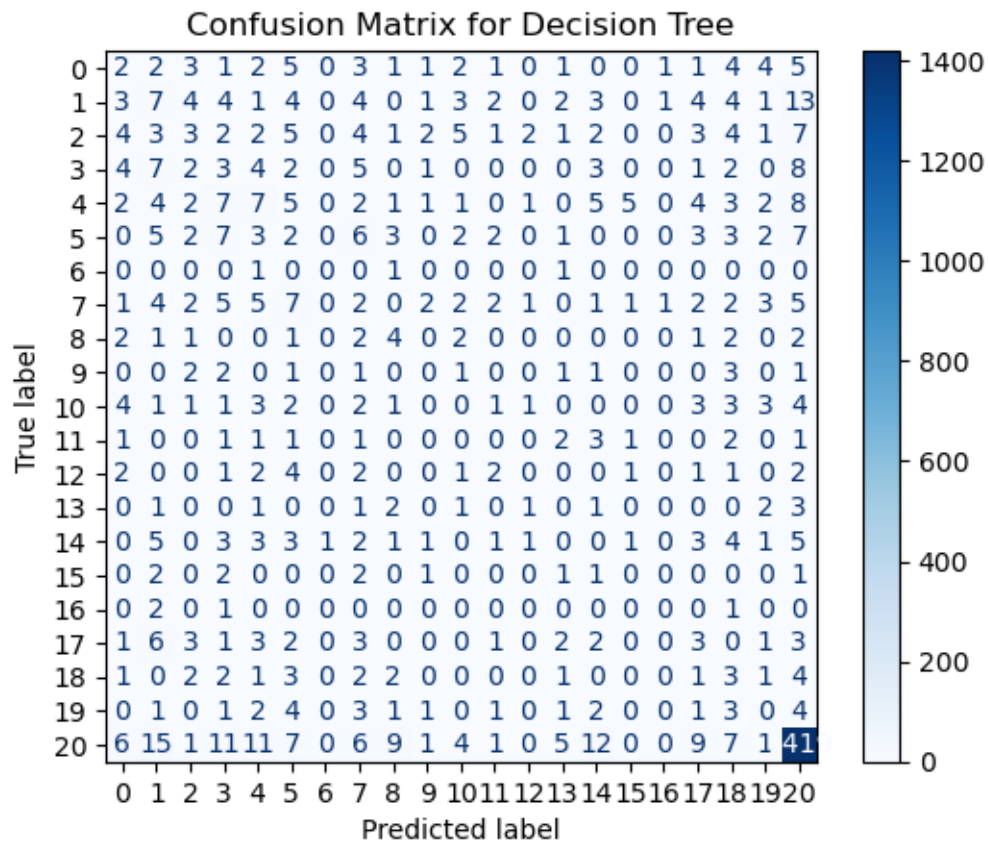
    confusion = confusion_matrix(y_test, y_pred_test)

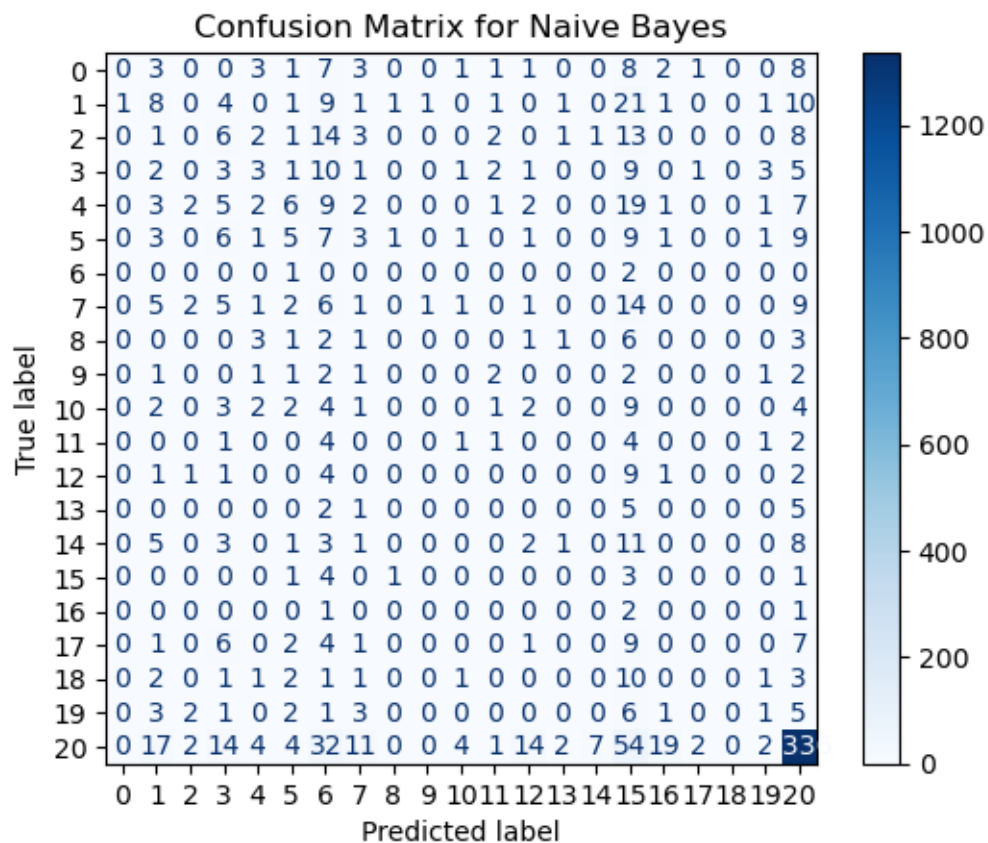
    results[name] = {
        "train_accuracy": train_accuracy,
        "test_accuracy": test_accuracy,
        "confusion_matrix": confusion
    }

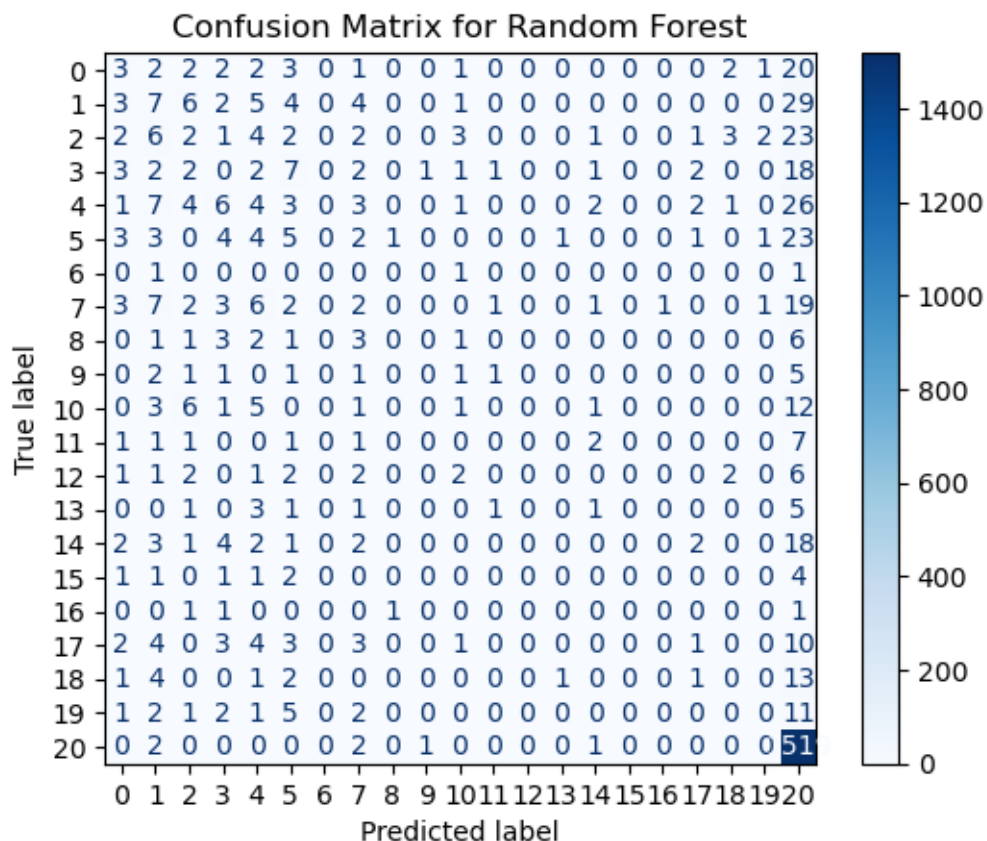
    disp = ConfusionMatrixDisplay(confusion_matrix=confusion)
    disp.plot(cmap=plt.cm.Blues)
    plt.title(f'Confusion Matrix for {name}')
    plt.show()

for name, metrics in results.items():
    print(f"{name}:")
    print(f"  Train Accuracy: {metrics['train_accuracy']:.4f}")
    print(f"  Test Accuracy: {metrics['test_accuracy']:.4f}")
    print("-" * 40)

```







Decision Tree:

Train Accuracy: 1.0000

Test Accuracy: 0.6886

Naive Bayes:

Train Accuracy: 0.6789

Test Accuracy: 0.6436

Random Forest:

Train Accuracy: 1.0000

Test Accuracy: 0.7307

Testing predictive models on new customers info, and getting prediction for churn and their possible churn reasons

```
[551]: import random

required_columns = [
    'Gender', 'Senior Citizen', 'Partner', 'Dependents', 'Phone Service',
    'Multiple Lines', 'Internet Service', 'Online Security', 'Online Backup',
```

```

    'Device Protection', 'Tech Support', 'Streaming TV', 'Streaming Movies',
    'Contract', 'Paperless Billing', 'Payment Method', 'CLTV', 'Churn Score',
    'Latitude', 'Longitude', 'Monthly Charges', 'Total Charges'
]

random_data = []
for i in range(7):
    tenure_months = random.randint(1, 72)
    random_customer = {
        'Gender': random.choice([0, 1]), # Assuming 0 for Female, 1 for Male
        'Senior Citizen': random.choice([0, 1]),
        'Partner': random.choice([0, 1]),
        'Dependents': random.choice([0, 1]),
        'Phone Service': random.choice([0, 1]),
        'Multiple Lines': random.choice([0, 1]),
        'Internet Service': random.choice([0, 1, 2]), # 0: DSL, 1: Fiber,
        'Optic': 2, # No service
        'Online Security': random.choice([0, 1]),
        'Online Backup': random.choice([0, 1]),
        'Device Protection': random.choice([0, 1]),
        'Tech Support': random.choice([0, 1]),
        'Streaming TV': random.choice([0, 1]),
        'Streaming Movies': random.choice([0, 1]),
        'Contract': random.choice([0, 1, 2]), # 0: Month-to-month, 1: One
        'Year': 2, # Two year
        'Paperless Billing': random.choice([0, 1]),
        'Payment Method': random.choice([0, 1, 2, 3]), # 0: Electronic check,
        'Mailed check': 2, # Bank transfer, 3: Credit card
        'CLTV': round(random.uniform(100, 1000), 2), # Random Customer
        'Lifetime Value':
        'Churn Score': round(random.uniform(0, 100), 2), # Random Churn Score
        'Latitude': round(random.uniform(34.0, 42.0), 6), # Random Latitude
        'Longitude': round(random.uniform(-120.0, -75.0), 6), # Random
        'Longitude':
        'Monthly Charges': round(random.uniform(20, 150), 2),
        'Total Charges': round(random.uniform(100, 5000), 2),
        'Tenure Months': tenure_months,
        'Popular Tenure': tenure_months > 12,
    }
    random_data.append(random_customer)

new_customers = pd.DataFrame(random_data)

print("Before Encoding:")
print(new_customers)

```

Before Encoding:

	Gender	Senior Citizen	Partner	Dependents	Phone Service	Multiple Lines	\
0	1	1	1	1	1	1	
1	1	0	0	1	1	0	
2	0	1	0	0	1	0	
3	0	1	1	1	1	1	
4	0	1	1	0	1	0	
5	0	1	1	1	1	1	
6	1	0	1	1	1	1	

	Internet Service	Online Security	Online Backup	Device Protection	...	\
0	1	0	1	1	...	
1	1	0	0	0	...	
2	0	0	0	0	...	
3	0	1	1	0	...	
4	1	0	0	1	...	
5	2	1	1	0	...	
6	1	0	1	1	...	

	Paperless Billing	Payment Method	CLTV	Churn	Score	Latitude	\
0	1	2	697.19		39.94	39.357528	
1	0	1	756.91		84.45	34.365494	
2	1	3	225.48		3.99	41.981300	
3	0	1	656.14		73.31	38.710241	
4	1	3	979.17		85.81	41.757198	
5	1	3	224.51		9.81	38.267444	
6	1	0	955.25		80.63	36.327863	

	Longitude	Monthly Charges	Total Charges	Tenure Months	Popular	Tenure
0	-102.082935	117.19	3209.91	35		True
1	-97.563329	111.07	2826.83	27		True
2	-114.533241	116.32	4765.55	6		False
3	-111.223638	149.49	3326.46	16		True
4	-110.214410	49.54	484.30	38		True
5	-83.833012	112.99	501.52	52		True
6	-77.435941	60.28	4630.11	61		True

[7 rows x 24 columns]

```
[553]: categorical_columns = [  
    'Gender', 'Senior Citizen', 'Partner', 'Dependents', 'Phone Service',  
    'Multiple Lines', 'Internet Service', 'Online Security', 'Online Backup',  
    'Device Protection', 'Tech Support', 'Streaming TV', 'Streaming Movies',  
    'Contract', 'Paperless Billing', 'Payment Method'  
]  
  
label_encoder = LabelEncoder()
```



```

for column in categorical_columns:
    new_customers[column] = label_encoder.fit_transform(new_customers[column])

print("\nAfter Encoding:")
print(new_customers)

```

After Encoding:

	Gender	Senior Citizen	Partner	Dependents	Phone Service	Multiple Lines	\
0	1	1	1	1	0	1	
1	1	0	0	1	0	0	
2	0	1	0	0	0	0	
3	0	1	1	1	0	1	
4	0	1	1	0	0	0	
5	0	1	1	1	0	1	
6	1	0	1	1	0	1	

	Internet Service	Online Security	Online Backup	Device Protection	...	\
0	1	0	1	1	...	
1	1	0	0	0	...	
2	0	0	0	0	...	
3	0	1	1	0	...	
4	1	0	0	1	...	
5	2	1	1	0	...	
6	1	0	1	1	...	

	Paperless Billing	Payment Method	CLTV	Churn Score	Latitude	\
0	1	2	697.19	39.94	39.357528	
1	0	1	756.91	84.45	34.365494	
2	1	3	225.48	3.99	41.981300	
3	0	1	656.14	73.31	38.710241	
4	1	3	979.17	85.81	41.757198	
5	1	3	224.51	9.81	38.267444	
6	1	0	955.25	80.63	36.327863	

	Longitude	Monthly Charges	Total Charges	Tenure Months	Popular Tenure
0	-102.082935	117.19	3209.91	35	True
1	-97.563329	111.07	2826.83	27	True
2	-114.533241	116.32	4765.55	6	False
3	-111.223638	149.49	3326.46	16	True
4	-110.214410	49.54	484.30	38	True
5	-83.833012	112.99	501.52	52	True
6	-77.435941	60.28	4630.11	61	True

[7 rows x 24 columns]

```
[684]: # Reorder the new_customers DataFrame to match the required columns
#new_customers = new_customers[required_columns]
training_features = [
    'Latitude', 'Longitude', 'Gender', 'Senior Citizen', 'Partner',
    ↪ 'Dependents',
    'Tenure Months', 'Phone Service', 'Multiple Lines', 'Internet Service',
    'Online Security', 'Online Backup', 'Device Protection', 'Tech Support',
    'Streaming TV', 'Streaming Movies', 'Contract', 'Paperless Billing',
    'Payment Method', 'Monthly Charges', 'Total Charges', 'Churn Score',
    'CLTV', 'Popular Tenure',
]

new_customers = new_customers[training_features]

# Now you can predict with the model
y_pred_churn_label = model.predict(new_customers)

# Random churn reason predictions
random_churn_reasons = ['Service Quality', 'Price', 'Competitor', 'Other']
y_pred_churn_reason = [random.choice(random_churn_reasons) for i in
    ↪ range(len(new_customers))]

# Instead of directly assigning to columns, use .loc[] to modify values
new_customers.loc[:, 'Predicted Churn Label'] = y_pred_churn_label
new_customers.loc[:, 'Predicted Churn Reason'] = y_pred_churn_reason

# Display the results
print(new_customers[['Predicted Churn Label', 'Predicted Churn Reason']])
print(type(model))
```

```
Predicted Churn Label Predicted Churn Reason
0                No      Service Quality
1                Yes      Service Quality
2                No                Other
3                No      Competitor
4                Yes                Price
5                No      Service Quality
6                Yes                Other
<class 'sklearn.ensemble._forest.RandomForestClassifier'>
```

2.2.2 Conclusions

This project successfully analyzed customer churn data and built a predictive system using data from a fictional telecommunications company with 7,043 customers in California.

Through **exploratory data analysis** using correlation matrices, chi-square tests, and other statistical techniques, the project was able to address the original questions:

- What is the most popular reason for customers canceling their subscription?

The most common reasons include dissatisfaction with support staff, competition, low download speeds, and limited data availability.

- **When do cancellations most commonly occur?**

Most churn happens within the **first month** of service.

- **Are the correlations between these factors statistically significant?**

- **Tenure Months** shows strong correlation with **Total Charges** and **CLTV**.
- **Monthly Charges** also strongly correlate with **Total Charges**.
- **Total Charges** is highly correlated with **Tenure Months**, **Monthly Charges**, and **CLTV**.

- **How do monthly charges affect the overall duration of subscriptions?**

Although the effect is statistically significant, **monthly charges alone explain only a small portion of the variability** in subscription duration. Other features likely play a more influential role.

- **How does the type of subscription contract influence churn?**

Month-to-month contracts show the **highest churn rate**, while one- and two-year contracts are much more stable and have significantly lower churn rates.

- **How can churn be reduced?**

- Improve customer support through better training
- Offer higher download speeds and more data
- Update or replace outdated devices
- Introduce more self-service options via the company website

2.2.3 Machine Learning Results

In the extended part of the project, **predictive models** were developed to classify both:

1. **Whether a customer will churn**
2. **The predicted reason for their churn**

Three models were compared: **Decision Tree**, **Naive Bayes**, and **Random Forest**.

Among them, the **Random Forest classifier consistently achieved the highest accuracy** in both churn classification and churn reason prediction.

This addition transforms the project from a purely analytical task into a **practical, data-driven decision support tool** that telecom companies can use to anticipate churn and tailor their retention strategies based on predicted reasons.

Let me know if you want it formatted as a PDF or Markdown, or if you'd like to add charts/tables for results!

2.3 Further Research

Based on our exploratory data analysis and initial modeling, there are several promising directions for further research:

1. Developing a Predictive ML Model for Churn:

- **Objective:**

A machine learning model was successfully developed to predict not only the likelihood of customer churn but also the underlying reason for churn. Among the tested algorithms, Random Forest demonstrated the highest accuracy for both tasks.

- **Outcome:**

A robust and interpretable predictive system capable of providing dual insights: whether a customer is likely to churn and why. This dual-layered prediction enables proactive, targeted interventions tailored to specific reasons (e.g., service quality, price, competition).

2. Personalization and Customer Segmentation:

- **Goal:**

Use the insights from the predictive model to segment customers into different risk categories.

- **Approach:**

- Develop clusters of customers based on their predicted churn risk and demographic or behavioral attributes.
- Tailor marketing strategies, service offerings, and retention initiatives (e.g., special discounts, personalized customer support) for each segment.

- **Outcome:**

Enhance customer engagement and increase subscription duration by offering individualized services that match customer needs.

3. Continuous Model Improvement:

- **Data Enrichment:**

- Incorporate additional data sources such as customer feedback, interaction logs, and external market data.
- Regularly update the model with new data to capture evolving customer behavior.

- **Operational Integration:**

- Deploy the model in a real-time environment to continuously monitor churn risk.
- Set up A/B tests to evaluate the impact of personalized interventions on customer retention.

- **Outcome:**

A dynamic, continuously improving system that adapts to changes in customer behavior and market conditions.

By pursuing these research directions, the company can not only predict customer churn with greater accuracy but also gain actionable insights into how to improve customer retention. This holistic approach can lead to more personalized service offerings, improved customer satisfaction, and ultimately, increased customer lifetime value.

Made by Alnur Nurumov

```
[705]: !jupyter nbconvert --to pdf "customer-churn-IBM-dataset(Predictive Model).ipynb"
```

```
[NbConvertApp] Converting notebook customer-churn-IBM-dataset(Predictive
```

```
Model).ipynb to pdf
[NbConvertApp] Support files will be in customer-churn-IBM-dataset(Predictive
Model)_files/
[NbConvertApp] Making directory ./customer-churn-IBM-dataset(Predictive
Model)_files
[NbConvertApp] Writing 160786 bytes to notebook.tex
[NbConvertApp] Building PDF
[NbConvertApp] Running xelatex 3 times: ['xelatex', 'notebook.tex', '-quiet']
[NbConvertApp] Running bibtex 1 time: ['bibtex', 'notebook']
[NbConvertApp] WARNING | bibtex had problems, most likely because there were no
citations
[NbConvertApp] PDF successfully created
[NbConvertApp] Writing 1517097 bytes to customer-churn-IBM-dataset(Predictive
Model).pdf
```

[]: